

# CONNECTTEL CUSTOMER CHURN PREDICTION

## PROJECT REPORT AND SUMMARY

PROJECT TITLE: ConnectTel Customer Churn Prediction using Supervised Machine Learning

AUTHOR: Adewale Odetara

DATE: 14th November, 2023

### Introduction:

In the dynamic landscape of telecommunications, customer churn poses a significant challenge for companies like ConnectTel. The ability to predict and understand customer churn is crucial for business sustainability and growth. In this project, I delve into the realm of churn prediction, leveraging machine learning techniques to develop models capable of identifying customers at risk of leaving the service.

### Project Background

ConnectTel is facing a client retention difficulty that threatens the company's long-term viability and growth. Customer churn prediction predicts potential customers to leave a company's service, requiring effective marketing strategies to increase their likelihood of staying.

### Project Objective

The primary goal is to develop an accurate and reliable predictive model using machine learning to predict which customers are likely to churn and implement proactive measures.

```
In [1]: # Import necessary Libraries
```

```
# For data analysis
import pandas as pd
import numpy as np
```

```
In [2]: # For data visualization
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [3]: # Data pre-processing
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler
from sklearn.preprocessing import LabelEncoder
```

```
In [4]: #Classifier Libraries
from sklearn.linear_model import SGDClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression
```

```
In [5]: # Ipip install xgboost
from xgboost import XGBClassifier
from sklearn.svm import LinearSVC, SVC
from sklearn.naive_bayes import GaussianNB
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier
```

```
In [6]: # Evaluation metrics
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, roc_auc_score
from sklearn.metrics import confusion_matrix
```

```
In [7]: import warnings
warnings.filterwarnings("ignore")
```

## Load the data

```
In [8]: # Load the dataset
df = pd.read_csv(r"C:\Users\ADMIN\Desktop\Resources\10Alytics Data Science\Capstone Project\Custome")
```

```
In [9]: df.head()
```

```
Out[9]:
```

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService
0	7590-VHVEG	Female	0	Yes	No	1	No	No phone service	DSL
1	5575-GNVDE	Male	0	No	No	34	Yes	No	DSL
2	3668-QPYBK	Male	0	No	No	2	Yes	No	DSL
3	7795-CFOCW	Male	0	No	No	45	No	No phone service	DSL
4	9237-HQITU	Female	0	No	No	2	Yes	No	Fiber optic

5 rows × 21 columns

```
In [10]: df.shape
```

```
Out[10]: (7043, 21)
```

## The data has 7043 rows and 21 columns

```
In [11]: # Data verification - Data type, number of features and rows, missing data, e.t.c
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   customerID            7043 non-null   object
1   gender                7043 non-null   object
2   SeniorCitizen         7043 non-null   int64
3   Partner              7043 non-null   object
4   Dependents            7043 non-null   object
5   tenure               7043 non-null   int64
6   PhoneService         7043 non-null   object
7   MultipleLines         7043 non-null   object
8   InternetService      7043 non-null   object
9   OnlineSecurity       7043 non-null   object
10  OnlineBackup          7043 non-null   object
11  DeviceProtection     7043 non-null   object
12  TechSupport          7043 non-null   object
13  StreamingTV          7043 non-null   object
14  StreamingMovies       7043 non-null   object
15  Contract              7043 non-null   object
16  PaperlessBilling     7043 non-null   object
17  PaymentMethod        7043 non-null   object
18  MonthlyCharges        7043 non-null   float64
19  TotalCharges         7043 non-null   object
20  Churn                7043 non-null   object
dtypes: float64(1), int64(2), object(18)
memory usage: 1.1+ MB

```

There is 21 columns in the dataset and only columns are numeric data type while the remaining are categorical data type.

```

In [12]: # Statistical Analysis of the data
df.describe()

```

```

Out[12]:

```

	SeniorCitizen	tenure	MonthlyCharges
count	7043.000000	7043.000000	7043.000000
mean	0.162147	32.371149	64.761692
std	0.368612	24.559481	30.090047
min	0.000000	0.000000	18.250000
25%	0.000000	9.000000	35.500000
50%	0.000000	29.000000	70.350000
75%	0.000000	55.000000	89.850000
max	1.000000	72.000000	118.750000

To understand the distribution of variables and the relationship between them.

```

In [13]: df.describe(exclude=["int64", "float64"]).T

```

Out[13]:

	count	unique	top	freq
<b>customerID</b>	7043	7043	7590-VHVEG	1
<b>gender</b>	7043	2	Male	3555
<b>Partner</b>	7043	2	No	3641
<b>Dependents</b>	7043	2	No	4933
<b>PhoneService</b>	7043	2	Yes	6361
<b>MultipleLines</b>	7043	3	No	3390
<b>InternetService</b>	7043	3	Fiber optic	3096
<b>OnlineSecurity</b>	7043	3	No	3498
<b>OnlineBackup</b>	7043	3	No	3088
<b>DeviceProtection</b>	7043	3	No	3095
<b>TechSupport</b>	7043	3	No	3473
<b>StreamingTV</b>	7043	3	No	2810
<b>StreamingMovies</b>	7043	3	No	2785
<b>Contract</b>	7043	3	Month-to-month	3875
<b>PaperlessBilling</b>	7043	2	Yes	4171
<b>PaymentMethod</b>	7043	4	Electronic check	2365
<b>TotalCharges</b>	7043	6531		11
<b>Churn</b>	7043	2	No	5174

In [14]: *# Check for duplicates*  
`df.duplicated().sum()`

Out[14]: 0


## Data Visualization

Investigate the data to discover any patterns.

I used seaborn countplot to plot a graph against churn column for the categorical data

In [15]: `sns.countplot(x='Churn', data=df, hue='gender')`

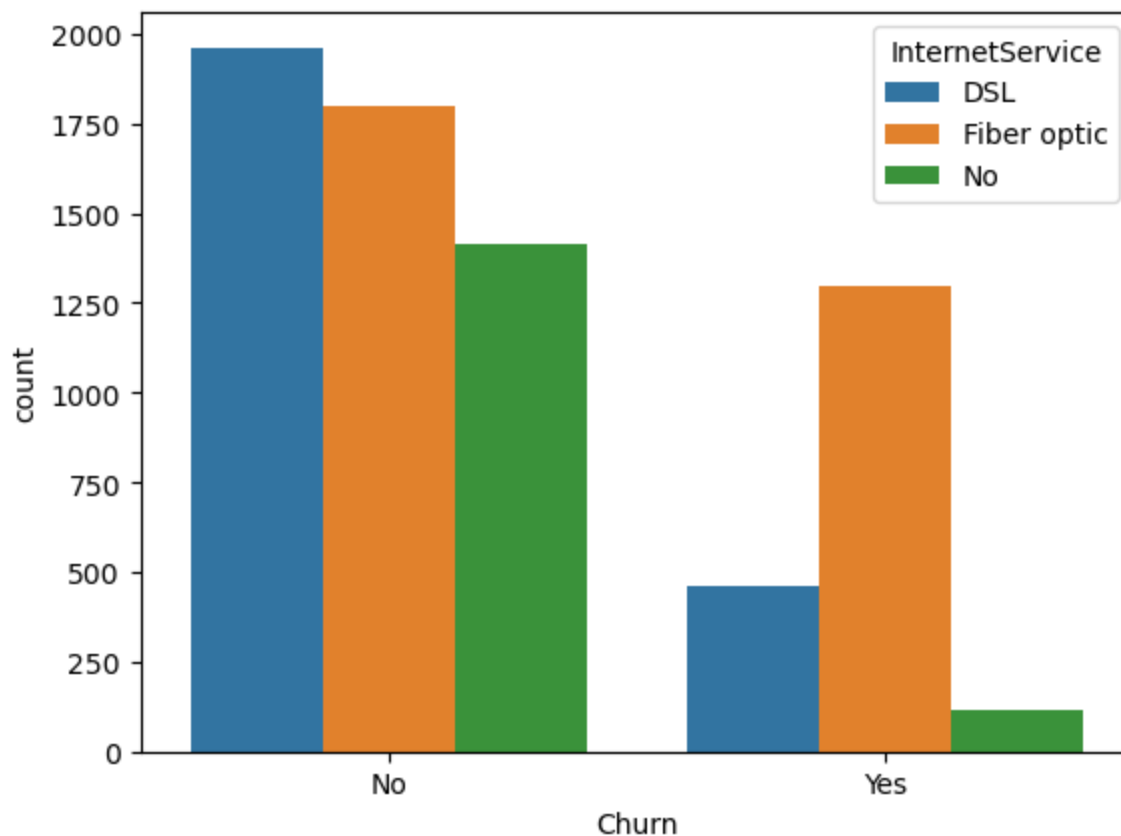
Out[15]: <Axes: xlabel='Churn', ylabel='count'>



The preceding plot shows that gender is not an important factor in customer churn in this data set because the numbers of both genders who have or have not churned are almost equal.

```
In [16]: sns.countplot(x='Churn',data=df, hue='InternetService')
```


```
Out[16]: <Axes: xlabel='Churn', ylabel='count'>
```



We can see that those who use fiber-optic services have a greater churn rate. This demonstrates that the company's Fiber-optic service has to be improved.

```
In [17]: sns.countplot(x='TechSupport',data=df, hue='Churn')
```

```
Out[17]: <Axes: xlabel='TechSupport', ylabel='count'>
```




Customers that do not have tech assistance have a higher turnover rate, which is obvious. This also demonstrates that the company's technical help is of high quality.

## Check for outliers


```
In [18]: sns.boxplot(x=df["tenure"])
```

```
Out[18]: <Axes: xlabel='tenure'>
```



```
In [19]: sns.boxplot(x=df["MonthlyCharges"])
```

```
Out[19]: <Axes: xlabel='MonthlyCharges'>
```



```
In [ ]:
```




# Exploratory Data Analysis


EDA is a comprehensive data analysis process that uses visual methods to uncover trends, patterns, and assumptions, while removing irregularities and unnecessary values from the data.

## Univariate Analysis

```
In [64]: plt.figure(figsize=(6,3))  
sns.countplot(x='Churn', data=df)  
plt.title('Distribution of Churn')  
plt.show()
```




```
In [21]: sns.histplot(x='MonthlyCharges', data=df, bins=30, kde=True)  
plt.title('Distribution of Monthly Charges')  
plt.show()
```



The company's success in retaining high-paying clients, even with monthly fees as high as \$100, is evident from the lack of clear patterns observed.

```
In [22]: sns.countplot(x='InternetService', data=df)
plt.title('Distribution of Internet Service Types')
plt.show()
```




In [ ]:


In [ ]:

## Bivariate Analysis

```
In [23]: # Bivariate analysis between MonthlyCharges and Churn
sns.boxplot(x='Churn', y='MonthlyCharges', data=df)
plt.show()
```



```
In [24]: sns.countplot(x='Contract', data=df, hue='Churn')  
plt.title('Churn by Contract Type')  
plt.show()
```

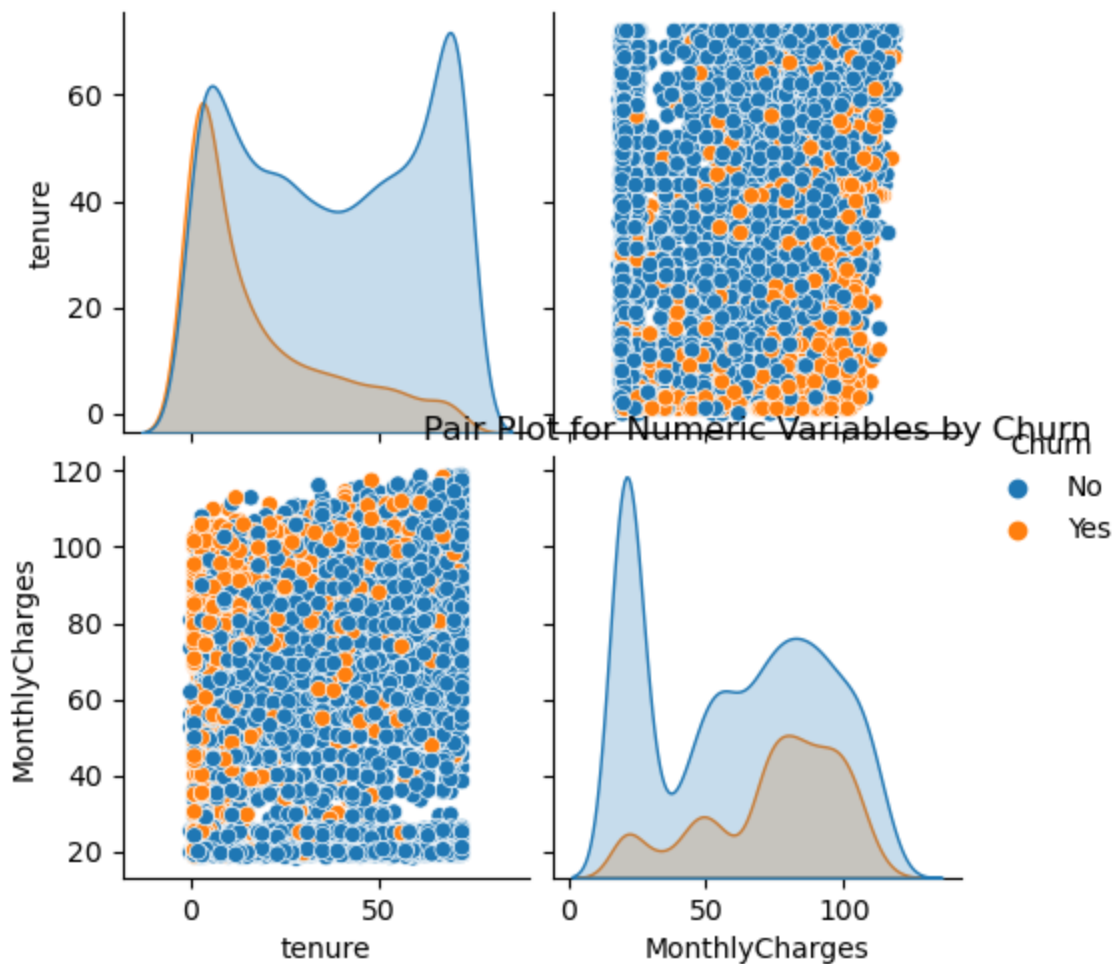


The churn rate is higher in the month-to-month, when new customers try out the service and decide whether to stay or terminate. This can be attributable primarily to the customer's uncertainty.


In [ ]:

## Multivariate Analysis

```
In [25]: sns.pairplot(df[['tenure', 'MonthlyCharges', 'TotalCharges', 'Churn']], hue='Churn')
plt.title('Pair Plot for Numeric Variables by Churn')
plt.show()
```



```
In [26]: corr_matrix = df.corr()
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm')
plt.title('Correlation Matrix')
plt.show()
```



In [ ]:

## Feature Engineering/Data Preprocessing

- Data Cleaning
- Encoding Categorical Variable
- Data Normalization

```
In [27]: # Create a copy of the data  
df1 = df.copy()
```

```
In [28]: df1.shape
```


```
Out[28]: (7043, 21)
```

```
In [29]: # Check for missing value  
df1.isnull().sum()
```

```
Out[29]: customerID      0
gender      0
SeniorCitizen  0
Partner      0
Dependents    0
tenure      0
PhoneService  0
MultipleLines  0
InternetService  0
OnlineSecurity  0
OnlineBackup  0
DeviceProtection  0
TechSupport   0
StreamingTV   0
StreamingMovies  0
Contract      0
PaperlessBilling  0
PaymentMethod  0
MonthlyCharges  0
TotalCharges  0
Churn         0
dtype: int64
```

```
In [30]: # Visualizing the missing data
plt.figure(figsize=(10,5))
sns.heatmap(df1.isnull(), cbar=True, cmap="Blues_r")
```

```
Out[30]: <Axes: >
```



## Encoding Categorical Variables

```
In [31]: # Encoding Categorical Variables

cat_feat = (df1.dtypes == "object")
cat_feat = list(cat_feat[cat_feat].index)

encoder = LabelEncoder()
for i in cat_feat:
    df1[i] = df1[[i]].apply(encoder.fit_transform)
```

```
In [32]: df1.head()
```

```
Out[32]:
```

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService
0	5375	0	0	1	0	1	0	1	0
1	3962	1	0	0	0	34	1	0	0
2	2564	1	0	0	0	2	1	0	0
3	5535	1	0	0	0	45	0	1	0
4	6511	0	0	0	0	2	1	0	1


5 rows × 21 columns

```
In [33]: # Drop CustomerID column
df1.drop('customerID', axis=1, inplace=True)
```

```
In [34]: # Explore Correlations

correlation_matrix = df1.corr()
plt.figure(figsize=(12, 10))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt=".2f")
plt.title('Correlation Matrix')
plt.show()
```





The correlation matrix shows how the variables are related to each other. a value close to 1 or -1 indicates a strong positive or negative correlation respectively. It can be seen that variable 'MonthlyCharges' and 'Paperlessbill' have a moderate positive correlation with outcome while 'tenue' and contract have a moderate negative relationship with the outcome, which indicating that they could be important factors in predicting customer churn.

```
In [35]: y = df1.pop('Churn')
```

```
In [36]: #df1.head()
```

## Create new feature from the dataset

```
In [37]: # Create a 'TotalTenureCharges' feature
df1['TotalTenureCharges'] = df1['tenure'] * df1['MonthlyCharges']
```

In [38]: `df1.head()`

Out[38]:

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineSecurit
0	0	0	1	0	1	0	1	0	
1	1	0	0	0	34	1	0	0	
2	1	0	0	0	2	1	0	0	
3	1	0	0	0	45	0	1	0	
4	0	0	0	0	2	1	0	1	

In [39]: `# Normalize/Scaling Dataset`

```
scaler = MinMaxScaler()
scaled_df = scaler.fit_transform(df1)
scaled_df = pd.DataFrame(scaled_df, columns=df1.columns)
```

In [40]: `scaled_df`

Out[40]:

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineS
0	0.0	0.0	1.0	0.0	0.013889	0.0	0.5	0.0	
1	1.0	0.0	0.0	0.0	0.472222	1.0	0.0	0.0	
2	1.0	0.0	0.0	0.0	0.027778	1.0	0.0	0.0	
3	1.0	0.0	0.0	0.0	0.625000	0.0	0.5	0.0	
4	0.0	0.0	0.0	0.0	0.027778	1.0	0.0	0.5	
...	...	...	...	...	...	...	...	...	...
7038	1.0	0.0	1.0	1.0	0.333333	1.0	1.0	0.0	
7039	0.0	0.0	1.0	1.0	1.000000	1.0	1.0	0.5	
7040	0.0	0.0	1.0	1.0	0.152778	0.0	0.5	0.0	
7041	1.0	1.0	1.0	0.0	0.055556	1.0	1.0	0.5	
7042	1.0	0.0	0.0	0.0	0.916667	1.0	0.0	0.5	

7043 rows × 20 columns

In [ ]:

## Build Machine Learning Model

In [41]: `# Split the dataset into training and testing sets - x = questions while y = answers`

```
X_train, X_test, y_train, y_test = train_test_split(scaled_df, y, test_size=0.2, random_state=42)
```

```
In [42]: # Model Building
# Logistic Regression

logreg = LogisticRegression()

logreg.fit(X_train, y_train)

ly_pred = logreg.predict(X_test)

print("Logistic Regression")
print("Accuracy:", accuracy_score(y_test, ly_pred))
print("Precision:", precision_score(y_test, ly_pred))
print("Recall:", recall_score(y_test, ly_pred))
print("F1-score:", f1_score(y_test, ly_pred))
print("AUC-ROC:", roc_auc_score(y_test, ly_pred))


Logistic Regression
Accuracy: 0.8176011355571328
Precision: 0.6858974358974359
Recall: 0.5737265415549598
F1-score: 0.6248175182481752
AUC-ROC: 0.7395659734801826
```

```
In [43]: # Create a confusion matrix

lcm = confusion_matrix(y_test, ly_pred)

# Visualize the confusion matrix

sns.heatmap(lcm, annot=True, cmap="Blues", fmt="g")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix")
plt.show()
```



```
In [44]: from sklearn.metrics import classification_report
```

```
In [45]: # Print the classification report - Logistic Regression

print(classification_report(y_test, ly_pred))
```

	precision	recall	f1-score	support
0	0.86	0.91	0.88	1036
1	0.69	0.57	0.62	373
accuracy			0.82	1409
macro avg	0.77	0.74	0.75	1409
weighted avg	0.81	0.82	0.81	1409

```
In [46]: # Model Building

# Random Forest Classifier

rfc = RandomForestClassifier()
rfc.fit(X_train, y_train)
rfy_pred = rfc.predict(X_test)
print("Random Forest")
print("Accuracy:", accuracy_score(y_test, rfy_pred))
print("Precision:", precision_score(y_test, rfy_pred))
print("Recall:", recall_score(y_test, rfy_pred))
print("F1-score:", f1_score(y_test, rfy_pred))
print("AUC-ROC:", roc_auc_score(y_test, rfy_pred))
```


Random Forest  
 Accuracy: 0.7906316536550745  
 Precision: 0.6413043478260869  
 Recall: 0.4745308310991957  
 F1-score: 0.5454545454545453  
 AUC-ROC: 0.6894854927696752

```
In [47]: # Create a confusion matrix

rcm = confusion_matrix(y_test, rfy_pred)

# Visualize the confusion matrix

sns.heatmap(rcm, annot=True, cmap="Blues", fmt="g")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix")
plt.show()
```



```
In [48]: # Print the classification report - Random Forest
print(classification_report(y_test, rfy_pred))
```

	precision	recall	f1-score	support
0	0.83	0.90	0.86	1036
1	0.64	0.47	0.55	373
accuracy			0.79	1409
macro avg	0.73	0.69	0.70	1409
weighted avg	0.78	0.79	0.78	1409

```
In [49]: # Model Building
```

```
# Decision Tree Classifier
```

```
dtc = DecisionTreeClassifier()
dtc.fit(X_train, y_train)
dty_pred = dtc.predict(X_test)
print("Decision Tree")
print("Accuracy:", accuracy_score(y_test, dty_pred))
print("Precision:", precision_score(y_test, dty_pred))
print("Recall:", recall_score(y_test, dty_pred))
print("F1-score:", f1_score(y_test, dty_pred))
print("AUC-ROC:", roc_auc_score(y_test, dty_pred))
```


```
Decision Tree
Accuracy: 0.723207948899929
Precision: 0.47780678851174935
Recall: 0.4906166219839142
F1-score: 0.4841269841269842
AUC-ROC: 0.6487832144668606
```

```
In [50]: # Create a confusion matrix
```

```
dcm = confusion_matrix(y_test, dty_pred)

# Visualize the confusion matrix

sns.heatmap(dcm, annot=True, cmap="Blues", fmt="g")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix")
plt.show()
```



```
In [51]: # Print the classification report
print(classification_report(y_test, dty_pred))
```

	precision	recall	f1-score	support
0	0.81	0.81	0.81	1036
1	0.48	0.49	0.48	373
accuracy			0.72	1409
macro avg	0.65	0.65	0.65	1409
weighted avg	0.73	0.72	0.72	1409

## Key Insights:

- Logistic Regression outperforms other models in terms of accuracy, precision, and recall.
- Decision Tree shows lower precision and recall compared to other models

In [52]: *# 7 Machine Learning Algorithms will be applied to the dataset*

```
classifiers = [[XGBClassifier(), 'XGB Classifier'],
               [RandomForestClassifier(), 'Random forest'],
               [SGDClassifier(), 'SGD Classifier'],
               [SVC(), 'SVC'],
               [GaussianNB(), "Naive Bayes"],
               [DecisionTreeClassifier(random_state = 42), "Decision tree"],
               [LogisticRegression(), 'Logistics Regression']]
```

In [53]: classifiers

Out[53]:

```
[[XGBClassifier(base_score=None, booster=None, callbacks=None,
               colsample_bylevel=None, colsample_bynode=None,
               colsample_bytree=None, device=None, early_stopping_rounds=None,
               enable_categorical=False, eval_metric=None, feature_types=None,
               gamma=None, grow_policy=None, importance_type=None,
               interaction_constraints=None, learning_rate=None, max_bin=None,
               max_cat_threshold=None, max_cat_to_onehot=None,
               max_delta_step=None, max_depth=None, max_leaves=None,
               min_child_weight=None, missing=nan, monotone_constraints=None,
               multi_strategy=None, n_estimators=None, n_jobs=None,
               num_parallel_tree=None, random_state=None, ...),
  'XGB Classifier'],
 [RandomForestClassifier(), 'Random forest'],
 [SGDClassifier(), 'SGD Classifier'],
 [SVC(), 'SVC'],
 [GaussianNB(), 'Naive Bayes'],
 [DecisionTreeClassifier(random_state=42), 'Decision tree'],
 [LogisticRegression(), 'Logistics Regression']]
```

In [54]:

```
acc_list = {}
precision_list = {}
recall_list = {}
roc_list = {}

for classifier in classifiers:
    model = classifier[0]
    model.fit(X_train, y_train)
    model_name = classifier[1]

    pred = model.predict(X_test)

    a_score = accuracy_score(y_test, pred)
```

```

p_score = precision_score(y_test, pred)
r_score = recall_score(y_test, pred)
roc_score = roc_auc_score(y_test, pred)

acc_list[model_name] = [str(round(a_score * 100, 2)) + '%']
precision_list[model_name] = [str(round(p_score * 100, 2)) + '%']
recall_list[model_name] = [str(round(r_score * 100, 2)) + '%']
roc_list[model_name] = [str(round(roc_score * 100, 2)) + '%']

if model_name != classifiers[-1][1]:
    print('')

```

```

In [55]: print("Accuracy Score")
s1 = pd.DataFrame(acc_list)
s1.head()

```

Accuracy Score

```

Out[55]:
   XGB Classifier  Random forest  SGD Classifier  SVC  Naive Bayes  Decision tree  Logistics Regression
0      77.86%      79.42%      79.49%  80.84%      75.94%      71.68%      81.76%

```

In [ ]:

```

In [56]: print("Precision Score")
s2 = pd.DataFrame(precision_list)
s2.head()

```

Precision Score

```

Out[56]:
   XGB Classifier  Random forest  SGD Classifier  SVC  Naive Bayes  Decision tree  Logistics Regression
0      59.44%      64.98%      73.08%  69.29%      53.11%      46.67%      68.59%

```

In [ ]:

```

In [57]: print("Recall Score")
s3 = pd.DataFrame(recall_list)
s3.head()

```

Recall Score

```

Out[57]:
   XGB Classifier  Random forest  SGD Classifier  SVC  Naive Bayes  Decision tree  Logistics Regression
0      51.47%      48.26%      35.66%  49.6%      77.75%      48.79%      57.37%

```

In [ ]:

```

In [58]: print("ROC Score")
s4 = pd.DataFrame(roc_list)
s4.head()

```

ROC Score



Out[58]:

	XGB Classifier	Random forest	SGD Classifier	SVC	Naive Bayes	Decision tree	Logistics Regression
0	69.41%	69.45%	65.46%	70.84%	76.52%	64.36%	73.96%

## Primary Metrics for Churn Prediction:

The two primary metrics for churn prediction are:

- Precision: Focuses on minimizing false positives, ensuring that customers predicted to churn are likely to do so.
- Recall: Emphasizes minimizing false negatives, ensuring that actual churners are correctly identified.

In [ ]:

## PROJECT REPORT AND SUMMARY

PROJECT TITLE: ConnectTel Customer Churn Prediction using Supervised Machine Learning

AUTHOR: Adewale Odetara

DATE: 14th November, 2023

### Introduction:

In the dynamic landscape of telecommunications, customer churn poses a significant challenge for companies like ConnectTel. The ability to predict and understand customer churn is crucial for business sustainability and growth. In this project, I delve into the realm of churn prediction, leveraging machine learning techniques to develop models capable of identifying customers at risk of leaving the service.

### Project Background

ConnectTel is facing a client retention difficulty that threatens the company's long-term viability and growth. Customer churn prediction predicts potential customers to leave a company's service, requiring effective marketing strategies to increase their likelihood of staying.

### Project Objective

The primary goal is to develop an accurate and reliable predictive model using machine learning to predict which customers are likely to churn and implement proactive measures.

### Data Loading and Cleaning

The project commenced with the crucial phase of data loading and cleaning, which involved:

- Preview the dataset to familiarize myself with its structure and contents.
- Identified and standardized data types for consistency.
- Detected and eliminated duplicate entries to ensure data integrity.

Exploratory Data Analysis (EDA) To uncover trends or patterns, obtain insights, and remove unnecessary values from the data, I conducted univariate, bivariate, and multivariate analyses to gain an in-depth

knowledge of the data and learn about its various features.

## Data Preprocessing:

1. Feature Engineering: • Identified and created relevant features that can contribute to the prediction of churn. • Handled missing values and outliers appropriately.
2. Encoding: • I used label encoding to convert categorical variables into a format suitable for machine learning models.
3. Scaling: • I normalized numerical features to ensure a level playing field for machine learning algorithms.

## Model Building:

Split the dataset into a training set and a testing set. A common split is 80% for training and 20% for testing.

- Dataset Split: The dataset was divided into 80% for training and 20% for testing.
- Models Implemented:
  - Logistic Regression
  - Random Forest
  - Decision Tree

## Model Evaluation:

The performance of each model was evaluated using key metrics: • Accuracy: overall correctness of the model predictions. • Precision: proportion of true positives among instances predicted as positive. • Recall: proportion of true positives among actual positive instances. • AUC (Area Under the Curve): the area under the ROC curve, measuring the model's ability to distinguish between classes.

## Key Insights:

- Logistic Regression stands out with the highest accuracy and a balanced precision-recall trade-off.
- Decision Tree shows lower precision and recall compared to other models.
- Random Forest provides a decent accuracy but has lower precision and recall compared to Logistic Regression. It may benefit from tuning hyperparameters to improve performance.
- Naïve Bayes is notable for high recall, making it suitable for scenarios where capturing all churn instances is crucial.
- Consider business priorities and the cost associated with false positives and false negatives when choosing a model.

## Confusion Matrix Analysis

### Interpretation:

- True Positives (TP): The number of customers correctly predicted as churners.
- True Negatives (TN): The number of customers correctly predicted as non-churners.
- False Positives (FP): The number of customers incorrectly predicted as churners.
- False Negatives (FN): The number of customers incorrectly predicted as non-churners.

## Model Comparison:

- Logistic Regression has the highest True Positives (214) and the fewest False Positives (159), indicating better performance in identifying actual churners.
- Random Forest has slightly fewer True Positives but also fewer False Positives compared to Decision Tree.
- Decision Tree has the highest False Positives and False Negatives, suggesting lower precision and recall compared to the other models.

## Recommendation:

- Logistic Regression appears to be performing better in this scenario, but the choice of the model depends on the specific goals and requirements of the business. Consider the trade-off between false positives and false negatives based on the business impact of predicting churn incorrectly.

The Logistic Regression model has a balanced distribution of false positives and false negatives. It demonstrates a higher ability to correctly identify non-churn instances (True Negatives) compared to correctly identifying churn instances (True Positives). The model's precision and recall can be further optimized by adjusting the classification threshold.

Similar to Logistic Regression, Random Forest exhibits a balanced distribution of false positives and false negatives. It performs slightly worse in correctly identifying both churn and non-churn instances compared to Logistic Regression. Random Forest's strength lies in ensemble learning, providing robustness against overfitting.

The Decision Tree model demonstrates a higher rate of false positives compared to both Logistic Regression and Random Forest. It shows comparable performance in correctly identifying churn instances (True Positives) but struggles with precision due to a higher false positive count. Decision Trees may benefit from pruning or tuning hyperparameters to improve overall performance.

## General Observations:

The models generally perform better at identifying non-churn instances (True Negatives) than identifying churn instances (True Positives). The choice between models depends on the specific business requirements and the importance of precision and recall in the context of customer churn.

## Conclusion and Recommendations:

In conclusion, the project highlights the potential of machine learning in predicting customer churn. ConnectTel can benefit from deploying the Logistic Regression model for its superior performance. Explore ensemble methods or model stacking to combine the strengths of different models.

Recommendations include continuous monitoring, periodic model updates, and leveraging insights from misclassifications to enhance model robustness and business strategies. Fine-tune model parameters and thresholds to balance precision and recall based on business goals. By prioritizing precision and recall, ConnectTel can strategically address customer churn, fostering long-term customer relationships and business success.

